Due Date April 25 Noon:
Late submissions might not be graded

Final Submission Instructions for Final Fixed Project:
You will submit the following on moodle:

1) Your report (as a pdf)
2) Code (either as a standalone ipynb, .py, or zip file with multiple scripts) that can reproduce your final results for Challenge 1 and 2. This code should train and test on 50 seeds from the testbed
3) Appendix Code (optional) if you would like to illustrate the code of other methods you have used you may include snippets of the code in the appendix of your report

Challenge 1 Submission Details:

Submit the final result of training and eval as per the testbed reporting the average and standard deviation. This should be done on a minimum of 10 splits and include the variance. It should be easy to reproduce these results if running your final submission script.

Challenge 2 Submission Details:
The submission instructions are the same as in Challenge 1 except for:

a) If your method takes long to run you may reduce the number of seeds, but should try to use at least 5.
b) If your final solution requires a lot of setup and might be hard to reproduce make sure the code is easy to follow and do your best to assure reproducibility.

Speed and Implementation Difficulty
Your final report should also include a discussion and measurement of the training and testing time of your method (or methods) as compared to the baseline code. Particularly for Challenge 1. These measurements can be approximate, but should give some idea of the relative speed to the baseline code from the testbed when run on the same hardware. For example “my method is 2x slower than the baseline in the testbed” etc. Having a slower method is not necessarily bad depending on other factor.

Particularly for challenge 2 Include a discussion (e.g. 1 paragraph) in your report of the complexity of implementing the final approach or implementing all the main approaches tried. Imagine you were at a company and had the same problem describe the simplicity to get working or lack thereof of the different methods you tried. This can describe things like “hyperparameters are hard to tune” or “many external libraries need to be installed” etc..
Particular

Length of Report

The report does not have a specific length requirement, it should describe what you tried and as well include the literature review. An approximate (soft) suggest is 3-10 pages depending on group size with two column 11 pt font (not including appendices with potential code snippets or other things you may want to include).

Suggested length of literature review is stated in the original project description as 1 page per work, this would be for two column 11 point font.

Relevance of literature review
If you are really interested in some work in deep learning that is distant from your best methods or this specific problem there is some leeway in having the literature review be on slightly different topics. For example you may do the literature review to discuss a paper about the architecture you have tried but didn't use in the end or about a method you wanted to try such as GANs but which was too cumbersome to implement or get working.

Alternative Project
Submission should include a clearly written report including literature review (at least 1 paper per grad student in project) and if sensible some *readable* code snippets. For report length follow guidelines as above (3-10 pages).